





Calibration and discrimination across exits

Dataset	Output	Accuracy (%)	ECE	NLL	Brier
IV-2a	S	83.02	0.0657	0.4921	0.2539
IV-2a	M	86.38	0.0802	0.4893	0.2291
IV-2a	D	86.57	0.0821	0.4902	0.2280
IV-2a	F	86.00	0.0822	0.4877	0.2286
IV-2a	Fusion	87.04	0.0780	0.4697	0.2232
IV-2b	S	86.83	0.0537	0.3115	0.1917
IV-2b	M	89.38	0.0521	0.2820	0.1677
IV-2b	D	89.05	0.0565	0.2852	0.1698
IV-2b	F	89.76	0.0522	0.2816	0.1650
IV-2b	Fusion	89.99	0.0531	0.2734	0.1624

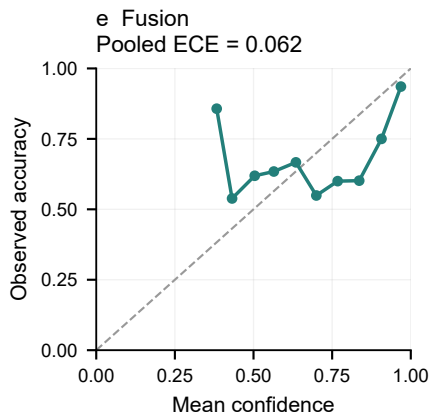
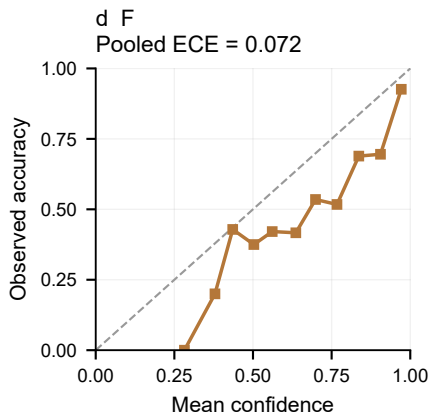
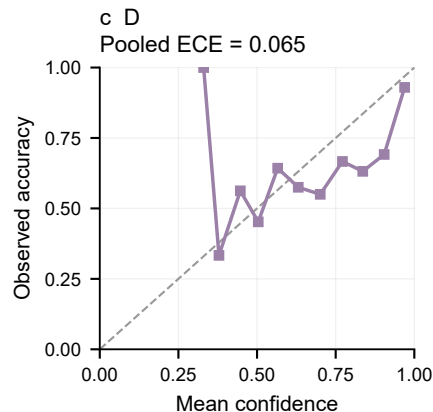
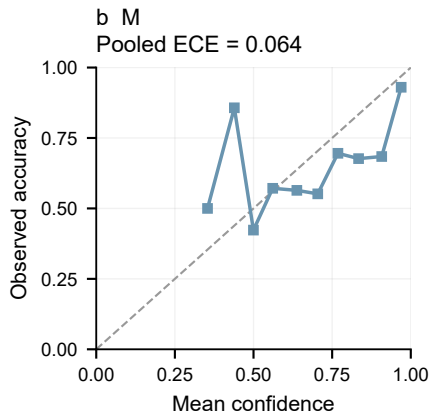
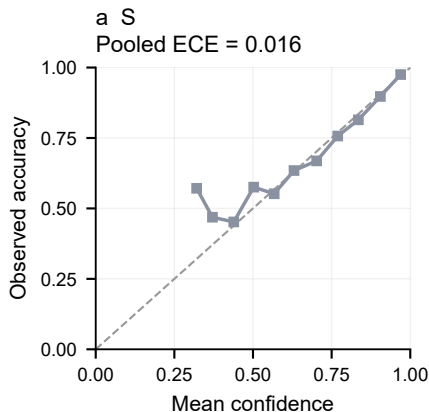
Equal-weight subject means; n = 9 per dataset. Lower ECE, NLL and Brier are better.
F is the final exit of the same multi-exit model, not an independently trained baseline.

Reliability of final-exit and fused predictions



All test trials pooled for visualization. Empty bins are not connected.
Calibration-table ECE is a subject mean, not the pooled-diagram ECE.

Exit-wise reliability | IV-2a



All outputs shown

S / M / D: auxiliary exits

F: final exit

Fusion: weighted logits

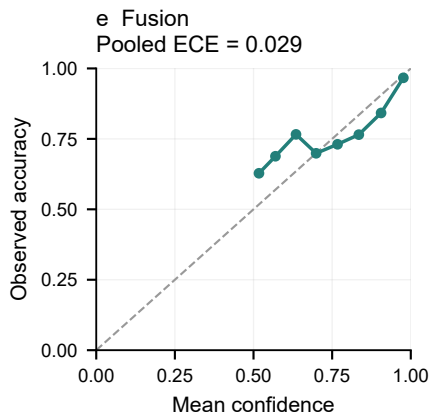
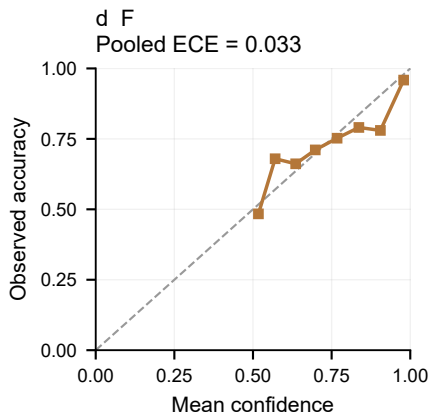
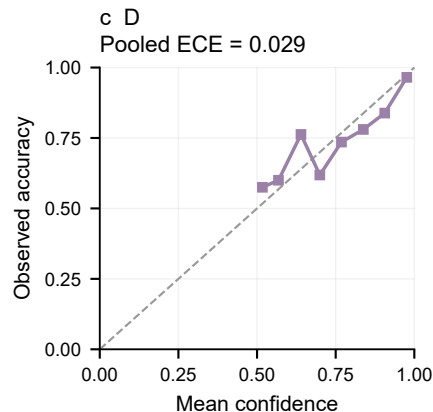
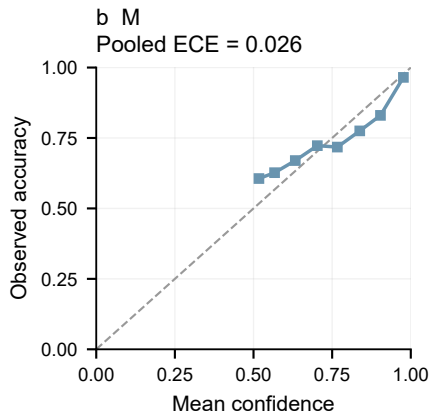
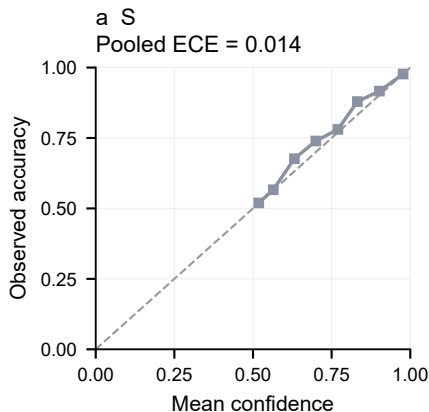
2,592 test trials

9 subjects

Table ECE uses subject means.
This figure shows pooled ECE.

Pooled test trials; 15 equal-width confidence bins. Unoccupied bins remain empty.

Exit-wise reliability | IV-2b



All outputs shown

S / M / D: auxiliary exits

F: final exit

Fusion: weighted logits

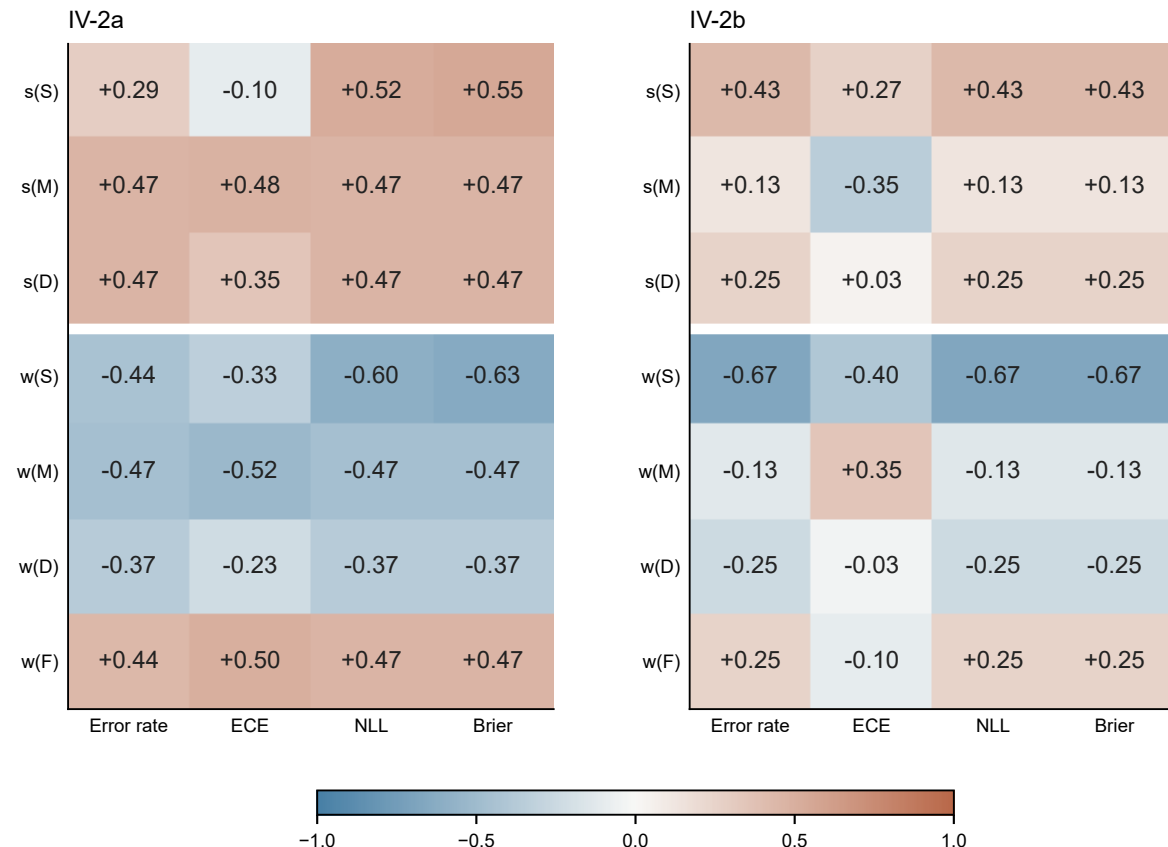
2,840 test trials

9 subjects

Table ECE uses subject means.
This figure shows pooled ECE.

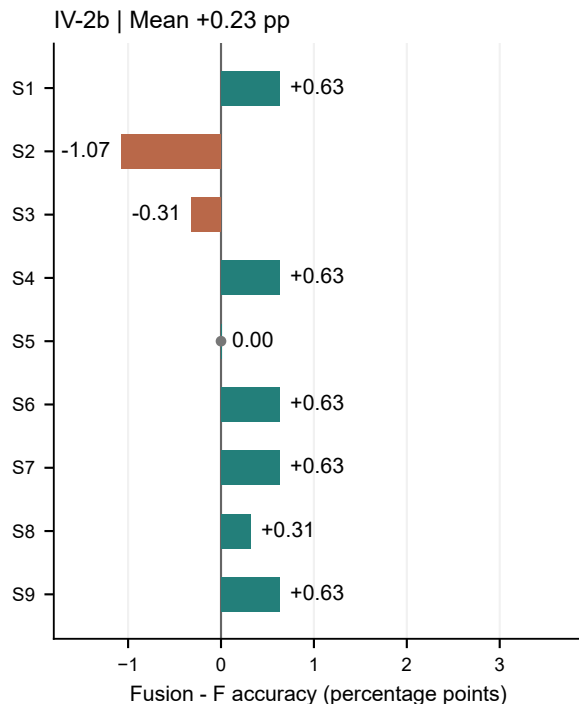
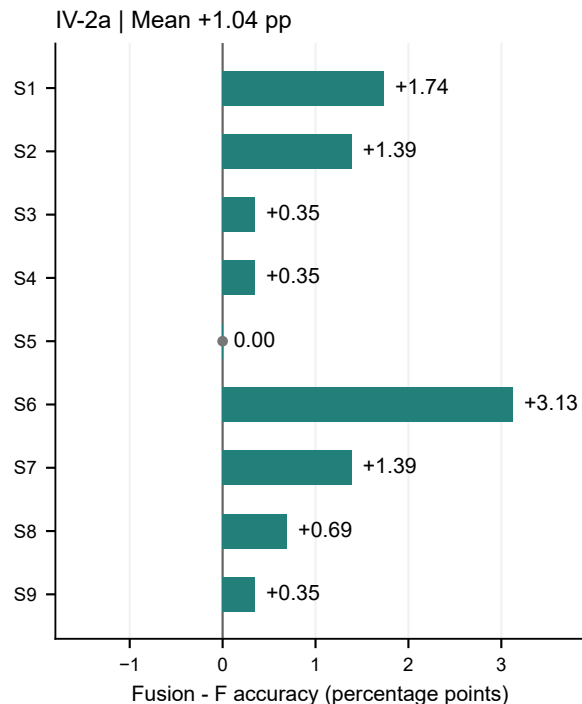
Pooled test trials; 15 equal-width confidence bins. Unoccupied bins remain empty.

Association of exit parameters with predictive reliability



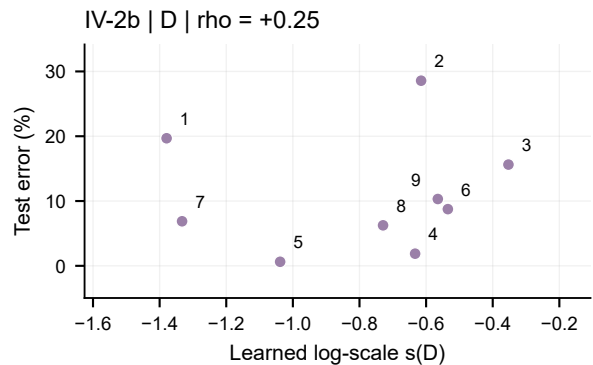
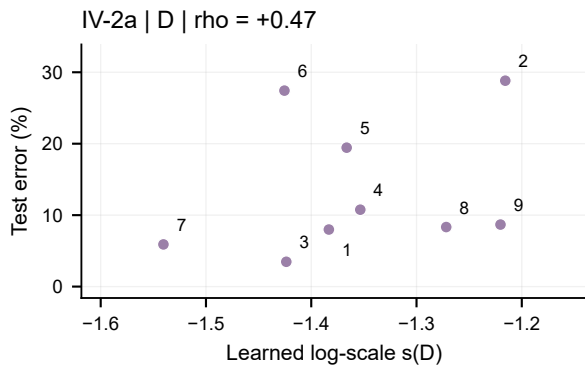
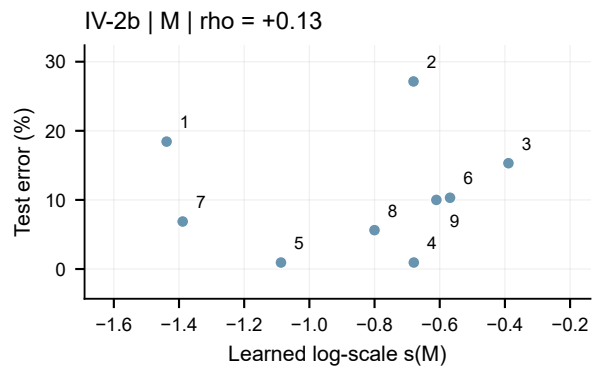
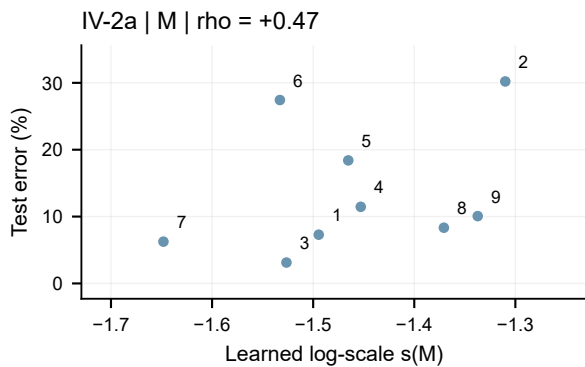
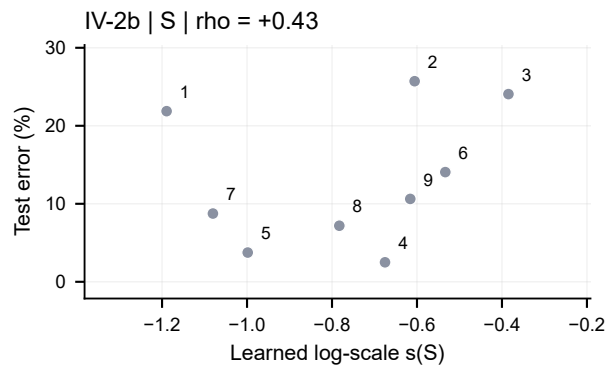
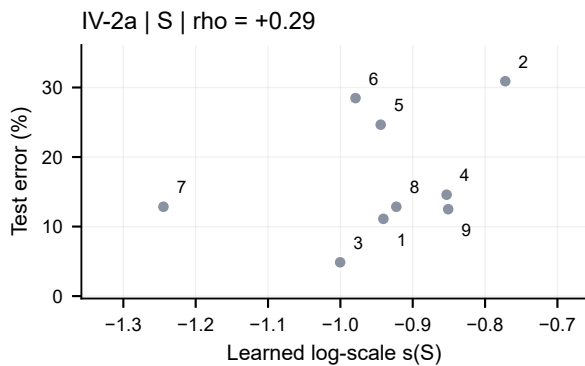
Spearman rho across 9 subjects, separately for each exit and dataset.
s: learned log-scale; w: normalized fusion weight. Lower column metrics are better.

Subject-level change from final exit to fusion



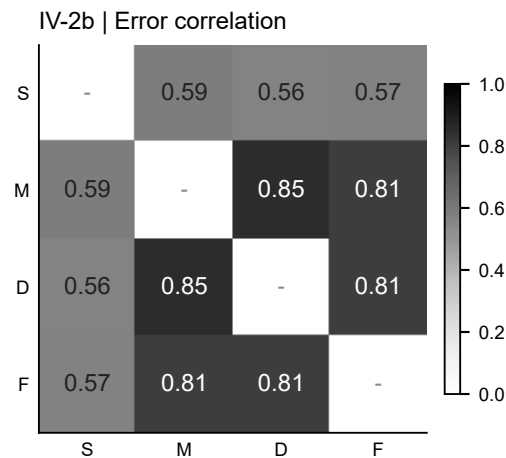
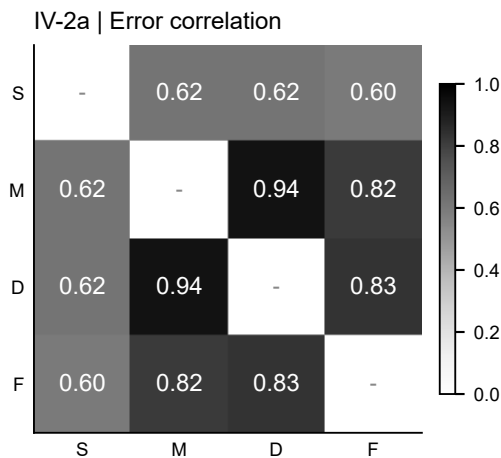
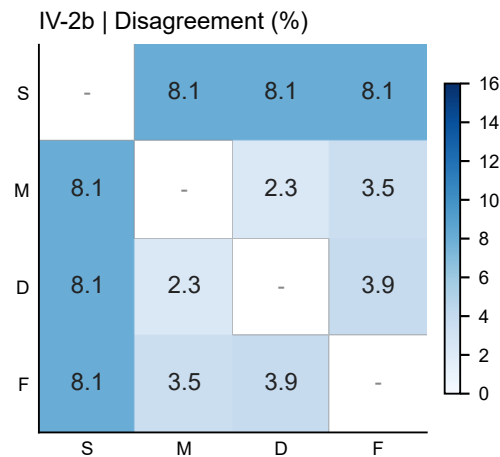
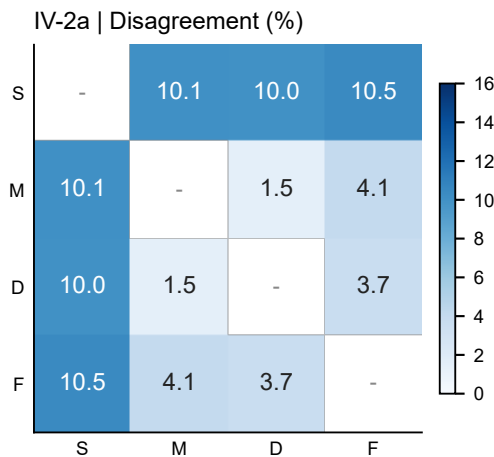
Each bar is one subject's paired accuracy change. Zero means unchanged accuracy.
The reference is F within the same multi-exit model, not an independent single-exit baseline.

Learned exit log-scales and test error rates



Every point is one subject; labels identify subjects 1-9. No fitted trend is imposed.
P values and the full 56-test correction are provided in the correlation table.

Prediction differences and shared errors across exits



Equal-weight subject means. Disagreement is not, by itself, evidence of useful complementarity. All six exit pairs and all subjects are included; diagonal cells are not comparisons.